



Arithmetic length generalization · RLT-1 4+4 vs 8+8 With Transformer baselines at depths 8 and 16 · seed 42

- RLT-1 4+4 (2,000)
- RLT-1 8+8† (1,200)
- Transformer 8 (2,000)
- Transformer 16 (2,000)



Same 256 test pairs per length · exact-answer bands: 95% Wilson intervals across pairs
Best ID-loss checkpoint (legend step); RLT-1 8+8† has reached 1,200/2,000 training steps.
RLT-1 4+4 and both Transformers completed 2,000 steps. Global batch 512, microbatch 32, width 512.