


Parity generalization · 16 layers



● RLT-1 8+8† (1400) ▲ RLT-1 10+6 (1800) ▼ RLT-1 12+4 (1000) ▲ RLT-1 14+2 (1500) ◆ RLT-1 16+0 (1900)
■ RLT-1 9+7 (1300) ◆ RLT-1 11+5 (1300) ◆ RLT-1 13+3 (2000) ▼ RLT-1 15+1 (1100) x Transformer 16 (2000)

Seed 42 · best ID-loss checkpoint in parentheses · 1,024 examples per length
† Training ongoing at snapshot: 8+8 at step 1,800